



# iSSDv2.0 RELEASE NOTES

---

Version 1.0

31<sup>st</sup> JULY, 2017

*Release Notes for iSSD v2.0*  
*(NVMe based intelligent SSD Card)*

## Revision History:

Date	Rev No.	Description	By
15-May-2017	1.0	Release Notes document for iSSDv2.0 (NVMe based intelligent SSD Card)	VVDN Technologies
31-July-2017	1.1	Fixed the crash issue and enhancement for iNIC Feature and NVMe	VVDN Technologies

## *Table of Contents*

<b>1</b>	<b>INTRODUCTION .....</b>	<b>4</b>
<b>2</b>	<b>RELEASE FEATURES.....</b>	<b>4</b>
2.1	BOARD SUPPORT PACKAGES.....	4
2.2	INIC .....	4
2.3	RoCE.....	4
2.4	NVME.....	5
<b>3</b>	<b>BUG FIXES, KNOWN ISSUES AND LIMITATIONS.....</b>	<b>5</b>
3.1	INIC .....	5
3.2	RoCE.....	6
3.3	NVME.....	9
<b>4</b>	<b>RELEASED ARTIFACTS .....</b>	<b>111</b>
4.1	RELEASE 1.0.....	111
4.2	RELEASE 1.1.....	11
<b>5</b>	<b>DOCUMENTS.....</b>	<b>122</b>

## 1 Introduction

This document provides release details for the LS2088a NVMe based iSSD project.

## 2 Release features

### 2.1 Board support packages

- Dual Boot partitions
- Firmware upgrade with binary tar ball
- Factory Default in bank 0
- BSP for iSSD( U-Boot, Linux and file system)
- Software versioning

#### 2.1.1 Release 1.0

- LS2088A default SDK is used
- Supports three operation mode as inic, issd, ls2-m2, are selectable at boot time. Stops at u-boot and 'setenvbootimgissd' and 'saveenv'.
- LS2 with M2.Card (ls2-m2) : NVMe host driver runs on ls2 and expose the M.2 card on it as nvme block device(/dev/nvme0n1). RoCE can use this nvme block device.

#### 2.1.2 Release 1.1

- iNIC Feature supported with Four SFP Ports.
- RoCE and NVMe Feature with FPGA DDR and NAND as storage.

### 2.2 iNIC

Detailed information about features implementations are elaborated in the document "*LS2085\_iNIC\_Feature\_Test\_Guide\_A0-01.pdf*". The implementation is same in LS2088 also.

### 2.3 RoCE

- Functionality test applications like "rping" are added.
- "qperf" is added for benchmarking.
- "ISER" (iSCSI Extensions for RDMA) is implemented for creating SAN (Storage Area Network).

**Note:** Detailed information about the test setup and the test procedures are elaborated in the document "*LS2\_iSSD\_RoCE\_implementation\_and\_testing\_guide\_A0-01.pdf*"

## 2.4 NVMe

- All mandatory Admin Commands are added.
- All mandatory NVM Commands are added.

### 2.4.1 Release 1.0

- LS2-M.2 Card is used as NVMe storage in ls2-m2 mode.
- FPGA DDR and NAND are used as storage in iSSD mode.

### 2.4.2 Release 1.1

- The inbound and outbound Feature in LS2 with x86 was implemented.
- The NVMe Feature was implemented for DDR and NAND as storage in FPGA.

## 3 Bug fixes, known issues and limitations

### 3.1 iNIC

Detailed information about bug fixes, known issues and limitations are elaborated in the document “*LS2085\_iNIC\_Feature\_Test\_Guide\_A0-01.pdf*”

#### 3.1.1 Release 1.0

- The iNIC feature cannot be tested with x86 machine since the PCIe outbound mapping is not working. Fix will be given in next release.
- iNIC feature between two LS2s can be used

#### 3.1.2 Release 1.1

- PCIe outbound Mapping was implemented in u-boot .
- iNIC feature is tested with x86 Machine.

##### 3.1.2.1 Bug fixes

1. Fixed crash issue due to socket buffer size is huge.

2. While transmitting larger packet size driver is not able to handle scatter gather packets from network layer

**Workaround:**

Added scatter gather support and sharing all the fragments information to LS2 instead of copying to single large packet.

3. While refilling buffers LS2 is not able to get buffers for rx packets.

**Workaround:**

Moved the refilling count from x86 to LS2 and maintaining refill ring for synchronization.

4. Not able to handle 8G memory of x86.

**Workaround:**

Mapping two outbound regions each of 4G and sharing 64 bit address to LS2.

**3.1.2.2 Enhancements**

- Added support for Four ports and large mtu size.
- Optimization made for performance improvement.

## 3.2 RoCE

Following are the bugs faced and their respective fixes

**Bug 1:**

When udp tunnel queuing fails then on the next run of `rx_e_arbiter` is run from `rx_e_arbiter_timer()` which was set sometime back in past and `list_empty(&rx_e->arbiter.qp_list)` will be true in this run, because `list_add_tail(qpl, &rx_e->arbiter.qp_list)` is not called on failure.

This halts scheduling of arbiter because of two reasons

1. Timer is not set from now, so `rx_e_arbiter_timer()` will not run again, so `rx_e_arbiter` is not scheduled
2. `arbiter_skb_queue` will schedule arbiter only if `arbiter_skb_count` is made to 1 from 0, which is not going to occur.

**File :** <kernel-directory>/drivers/infiniband/hw/rxe/rxe\_arbiter.c

**Function :** rxe\_arbiter()

**Fix 1 :**

1. Perform list\_add\_tail(qpl, &rxe->arbiter.qp\_list) which allows to make qp\_list non empty and arbiter will definitely attempt to queue skb to udp\_tunnel.
2. set timer which calls rxe\_arbiter\_timer() on timeout .

**Bug 2:**

Packet psn would be requestedpsn, but resource is not related to the requested packet, since resource not cleared for the request for which RNR\_NAK was sent.

**File :** <kernel-directory>/drivers/infiniband/hw/rxe/rxe\_resp.c

**Function :** read\_reply()

**Fix 2 :**

Adjusting code flow, all requests to bypasscheck\_psn() till Request for which RNR\_NAK was send is completed.

- Drop all the requests above the PSN for which RNR\_NAK was sent (No need as these request will come again).
- Allow Duplicate Packets to be processed (The requests before RNR\_NAK can be processed)
- Allow the subset of request for which RNR\_NAK was sent and clear sent\_psn\_rnr\_nak.
- Coded to prevent the formation of faulty headers, which occurred in default code
- Added fields in the structure rxe\_resp\_info for storing NAK details

**Bug 3:**

Retrial with wrong DMA LENGTH (when subset of earlier request is made, full DMA length was used)

**File :** <kernel-directory>/drivers/infiniband/hw/rxe/rxe\_req.c

**Function :** req\_retry()

**Fix 3 :**

Reduce the DMA LENGTH for the subset of request

**Expected errors:**

When the packets are dropped by any of the NIC, then we cannot recover the RC (Reliable Connection) connection beyond a point ie., if Resource for the Request is cleared and a duplicate request comes for the cleared resource, then the request is silently dropped. This may lead to timeout and closing the RC connection.

check "ifconfig" for the NIC card, and verify the fields errors,dropped, overruns for both TX and RX

**3.2.1.1 Bugs fixed**

1. Tx. enqueue error occurs in transmission side.

**Fix:**

Sysfs entries to select the control method and value for the respective method is added, please refer the document *LS2\_iSSD\_RoCE\_implementation\_user\_guide\_A0-02.docx*(section 8.3.4)



### 3.2.1.2 Known issue

#### 1. Out of order packets

**Reason:**

Multiple Queues in linuxQdisc can cause out of order sending of packets to DPAA transmit (hard transmit func). Receiving side reordering has to be implemented. For back to back connection `tc` command can be used to set one transmit queue which reduces out of order.

#### 2. iSERnoop timeout

**Reason:**

In file `/etc/iscsi/iscsid.conf`, set the below parameters

```
node.conn[0].timeo.noop_out_interval = 120
node.conn[0].timeo.noop_out_timeout = 240
```

which will enable more stable iSER connection, preventing iSER ping timeout which may occur when lot of traffic is going on iSER.

If iscsid is already running before you change the file `iscsid.conf` then

1. unmount and logout all iser sessions
2. killalliscsid
3. edit file as mentioned above
4. perform discovery and login

## 3.3 NVMe

Known issues and their workarounds are as follows:

### 3.3.1 Release 1.0

The Inbound and Outbound feature in LS2 is not working with x86 machine. So the NVMe feature cannot be tested now. Fix will be given in next release.

### 3.3.2 Release 1.1

- The Inbound and outbound feature in LS2 was fixed.
- The NVMe feature was tested.

---

### 3.3.3 Known issues and Bug fixes

#### 3.3.3.1 Bug Fixes

1. Host system BIOS failed to assign address to PCI BARs when PCI BAR Configurations are done in U-boot.

**Fix:**

Moved the PCI3 BAR Configuration to PBL from U-boot.

**NOTE:** Host system detects the card as two PCI functions(both iNIC and iSSD). But either iNIC or iSSD will work at a time based on *booting* env of u-boot

2. In the x86 system booting, at the time of nvme module (built in) loading crash happens and also not able to unload the nvme module.

**Workaround:**

Blacklist the nvme module and load the nvme module (nvme host driver) shared. The instructions are mentioned in the document *NVMe\_USER\_MANUAL.pdf*.

**Fix:**

Added default NVMe register values in FPGA itself after a reboot/power cycle, enabling the host to see proper NVMe register values for the host while NVMe probe happens.

#### 3.3.3.2 Known Issues

1. Could not map complete memory of x86 to LS2.

**Workaround:**

Map 4G memory of x86 to LS2 using the instructions from the document *NVMe\_USER\_MANUAL.pdf*

## 4 Released artifacts

### 4.1 Release 1.0

iSSD\_SW\_01.00.00

```
-- documents
| |-- LS2085_iNIC_Feature_Test_Guide_A0-00.pdf
| |-- LS2_iSSDv2.0_Release_notes_V-1.0.pdf
| |-- LS2_iSSD_RoCE_implementation_and_testing_guide_A0-01.pdf
| |-- LS2_iSSDv2.0_USER_MANUAL_A0-01.pdf
|
|-- issd-images_01.00.00.tar.gz
| |-- checksums.md5
| |-- flash_images.sh
| `-- images
|     |-- device-tree_01.00.00.dtb
|     |-- dpc_01.00.00.dtb
|     |-- dpl_01.00.00.dtb
|     |-- fpga_01.00.00.rpd
|     |-- mc_01.00.00.itb
|     |-- pbl_01.00.00.bin
|     |-- rootfs_01.00.00.sqsh
|     |-- u-boot_01.00.00.bin
|     |-- uimage_inic_01.00.00.bin
|     `-- uimage_issd_01.00.00.bin
|
|-- issd-sdk_01.00.00.iso
|-- README.txt
```

**Note:** Please refer user manual “LS2\_iSSDv2.0\_USER\_MANUAL\_A0-01.pdf” for hardware setup, software setup, flashing binaries into the target board and etc.

### 4.2 Release 1.1

iSSD\_SW\_01.00.01

```
-- documents
| |-- LS2085_iNIC_Feature_Test_Guide_A0-00.pdf
| |-- LS2_iSSDv2.0_Release_notes_V-2.0.pdf
| |-- LS2_iSSD_RoCE_implementation_and_testing_guide_A0-01.pdf
| |-- LS2_iSSDv2.0_USER_MANUAL_A0-01.pdf
|
|-- issd-images_01.00.00.tar.gz
| |-- checksums.md5
| |-- flash_images.sh
| |-- ls2088aissd
| `-- images
|     |-- device-tree_88_01.00.01.dtb
```

```
| -- dpc_88_01.00.01.dtb
| -- dpl_88_01.00.01.dtb
| -- FPGA_A3_03.01.00.rpd
| --FPGA_A7_03.03.00.rpd
| -- mc_88_01.00.01.itb
| -- pbl_88_01.00.01.bin
| -- rootfs_88_01.00.01.sqsh
| -- u-boot_88_01.00.01.bin
| -- uimage_inic_88_01.00.01.bin
| `-- uimage_issd_88_01.00.01.bin
|
|-- issd-sdk_01.00.00.iso
|-- README.txt
```

**Note:** For ls2-m2 use the uimage\_ls2\_m2\_88\_01.00.00.bin and rootfs\_ls22\_88\_01.00.00.bin images in m.2 directory in github. Please refer user manual “**LS2\_iSSDv2.0\_USER\_MANUAL\_A0-01.pdf**” for hardware setup, software setup, flashing binaries into the target board and etc.

## 5 Documents

File Name	Description
LS2_iSSDv2.0_Release_notes_V-1.1.pdf	Release notes
LS2_iSSDv2.0_USER_MANUAL_A0-01.pdf	User manual for LS2085a iSSD
LS2085_iNIC_Feature_Test_Guide_A0-01.pdf	This document describes the setup and the steps required to validate the Freescale LS2085 card iNIC functionality.
LS2_iSSD_RoCE_implementation_and_testing_guide_A0-03.pdf	This document describes the steps required for establishing and maintaining the RDMA
NVMe_USER_MANUAL_A00-02.pdf	NVMe user manual